



1 Overview

Welcome to HBase Machine Practice. It is highly recommended that you practice the **Tutorial: Introduction to HBase** before beginning this assignment. This MP has both java and python templates. Please use whichever you prefer.

2 General Requirements

Please note that our grader runs on a docker container and is **NOT connected to the internet**. Therefore, **no additional libraries are allowed** for this assignment. Also, you will **NOT be allowed to create any file or folder outside the current folder** (that is, you can only create files and folders in the folder that your solutions are in).

3 Set Up HBase

Step 1: Start the "default" Docker machine that you created when following the "Tutorial: Docker installation" in week 4, run:

```
1 docker-machine env
2 # follow the instruction to configure your shell: eval $(...)
3 docker-machine start default
```

Step 2: Download the Dockerfile and related files for this MP, change the current folder, build, and run the docker image, run:

```
1 git clone https://github.com/UIUC-public/Docker_MP3.git
2 cd Docker_MP3
3 docker build -t mp3 .
4 docker run -it mp3 bin/bash
```

Step 3: start HBase, run

```
1 start-hbase.sh
```

Java submission

****If you choose to do this assignment in Python, skip this part and go to "Python submission" part below.**

1 Requirements

This assignment will be graded based on **JDK 8**

2 Procedures

Step 1: Download the Java templates and change the current folder, run:

```
1 git clone https://github.com/UIUC-public/MP3_java.git
2 cd MP3_java
```

Step 2: Set the environment variable CLASSPATH

```
1 export CLASSPATH=$(hbase classpath)
2 export CLASSPATH="/MP3_java:${CLASSPATH}"
```

Step 3: Finish the exercises by editing the provided templates files. All you need to do is complete the parts marked with **TODO**. Please note that you are **NOT allowed to import any additional libraries**.



- Each exercise has a Java code template. All you must do is edit this file.
- Each exercise should be implemented in one file only. Multiple file implementation is not allowed.
- The code should be compiled and run on the sample Docker image.
- You should run the parts in the order described below. Otherwise, you might not get the correct results. For example, you should create the tables first and then list them.
- Remember that the output is case sensitive. Furthermore, your application should not remove the table and its values at the end of execution.

step 4: [updated]After you are done with the assignments, please check out the submission instructions under week 11 to submit your files..

Exercise A: Create HBase Tables

In this exercise, you will create 2 tables as shown below. To make the implementation easier, we have provided a boilerplate for this exercise in the following file: **TablePartA.java** .

Table 1: a table named "powers" in HBase with three column families as given below. For simplicity, treat all data as **String** (i.e., don't use any numerical data type).

| row key | personal | professional | custom |
|---------|----------|--------------|--------|
|---------|----------|--------------|--------|

Table 2: a table named "food" in HBase with two column families as given below. For simplicity, treat all data as **String** (i.e., don't use any numerical data type).

| row key | nutrition | taste |
|---------|-----------|-------|
|---------|-----------|-------|

Following are the commands we will use to run your file:

```
1 javac TablePartA.java
2 java TablePartA
```

Your output should contain:

```
INFO [main] client.HBaseAdmin: Created powers
INFO [main] client.HBaseAdmin: Created food
```

Exercise B: List HBase Tables

In this exercise, you will list all the tables created in the previous part. To make the implementation easier, we have provided a boilerplate for this exercise in the following file: **TablePartB.java** .

Following are the commands we will use to run your file:

```
1 javac TablePartB.java
2 java TablePartB
```

Your output should contain:

```
food
powers
```

Exercise C: Populate HBase Table with Data

In this exercise, you will insert data in the “powers” table created in the previous part A with following schema according to **input.csv**. You can assume that **input.csv** is in the same folder as the java files. To make the implementation easier, we have provided a boilerplate for this exercise in the following file: **TablePartC.java**.

| row key | personal | | professional | | custom |
|---------|----------|-------|--------------|----|--------|
| | hero | power | name | xp | color |
| row1 | | | | | |

Following are the commands we will use to run your file:

```
1 javac TablePartC.java
2 java TablePartC
```

Exercise D: Read Data

In this exercise, you will read some of the data in the populated “powers” table created in part C. To make the implementation easier, we have provided a boilerplate for this exercise in the following file: **TablePartD.java**.

Read the values for the following row ids in order and the given attributes:

Id: "row1", Values for (hero, power, name, xp, color)

Id: "row19", Values for (hero, color)

Id: "row1", Values for (hero, name, color)

Following is an example of output lines, please use spaces between the words. Your output should contain three lines for this exercise.

```
hero: no, color: green
```

Following are the commands we will use to run your file:

```
1 javac TablePartD.java
2 java TablePartD
```

Exercise E: Scan Data

In this exercise, you will scan the data in the populated “powers” table created in part C. To make the implementation easier, we have provided a boilerplate for this exercise in the following file: **TablePartE.java**.

Following are the commands we will use to run your file:

```
1 javac TablePartE.java
2 java TablePartE
```

The output should look similar to below:

```
keyvalues={row1/custom:color/1520644330258/Put/vlen=5/seqid=0, row1/personal:hero/1520644330258/Put/vlen=3/seqid=0, row1/personal:power/1520644330258/Put/vlen=3/seqid=0, row1/professional:name/1520644330258/Put/vlen=6/seqid=0, row1/professional:xp/1520644330258/Put/vlen=3/seqid=0}
keyvalues={row10/custom:color/1520644330298/Put/vlen=4/seqid=0, row10/personal:hero/1520644330298/Put/vlen=2/seqid=0, row10/personal:power/1520644330298/Put/vlen=5/seqid=0, row10/professional:name/1520644330298/Put/vlen=7/seqid=0, row10/professional:xp/1520644330298/Put/vlen=2/seqid=0}
keyvalues={row11/custom:color/1520644330301/Put/vlen=5/seqid=0, row11/personal:hero/1520644330301/Put/vlen=2/seqid=0, row11/personal:power/1520644330301/Put/vlen=3/seqid=0, row11/professional:name/1520644330301/Put/vlen=6/seqid=0, row11/professional:xp/1520644330301/Put/vlen=2/seqid=0}
```

Python submission

If you choose to do this assignment in Java, skip this part and go to "Java submission**" part above

1 Requirements

This assignment will be graded based on **Python 3**



2 Procedures

Step 1: Download the Python templates and change the current folder, run,

```
1 git clone https://github.com/UIUC-public/MP3_py.git
2 cd MP3_py
3
```

Step 2: Start Thrift, run

```
1 hbase thrift start &
```

note that you may want to press "enter" after below INFO to continue

```
INFO [main] thrift.ThriftServerRunner: starting TBoundedThreadPoolServer on /0.0.0.0:9090 with readTimeout 60000ms; min worker threads=16, max worker threads=1000, max queued requests=1000
```

step 3: Finish the exercises by editing the provided templates files. Please note that you are **NOT allowed to import any additional libraries**.

- Each exercise has one or more code template. All you must do is edit these files.
- The code will be run on the provided Docker image.

step 5: [updated]After you are done with the assignments, please check out the submission instructions under week 11 to submit your files.

Exercise A: Create HBase Tables

In this exercise, you will create 2 tables as shown below. To make the implementation easier, we have provided a boilerplate for this exercise in the following file: **TablePartA.py**.

Table 1: a table named "powers" in HBase with three column families as given below. For simplicity, treat all data as **String** (i.e., don't use any numerical data type).

| row key | personal | professional | custom |
|---------|----------|--------------|--------|
|---------|----------|--------------|--------|

Table 2: a table named "food" in HBase with two column families as given below. For simplicity, treat all data as **String** (i.e., don't use any numerical data type).

| row key | nutrition | taste |
|---------|-----------|-------|
|---------|-----------|-------|

We will run **python3 TablePartA.py**.

Your output should contain:

```
INFO [thrift-worker-0] client.HBaseAdmin: Created powers
INFO [thrift-worker-0] client.HBaseAdmin: Created food
```

Exercise B: List HBase Tables

In this exercise, you will list all the tables created in the previous part. To make the implementation easier, we have provided a boilerplate for this exercise in the following file: **TablePartB.py**. Command line and the expected output are as below

```
[root@2a4c5e229080:/MP3_py# python3 TablePartB.py
[b'food', b'powers']
```



Exercise C: Populate HBase Table with Data

coursera

In this exercise, you will insert data in the “powers” table created in the previous part A with following schema according to **input.csv**. You can assume that **input.csv** is in the same folder as the java files. To make the implementation easier, we have provided a boilerplate for this exercise in the following file: **TablePartC.py**.

| row key | personal | | professional | | custom |
|---------|----------|-------|--------------|----|--------|
| | hero | power | name | xp | color |
| row1 | | | | | |

We will run **python3 TablePartC.py**.

Exercise D: Read Data

In this exercise, you will read some of the data in the populated “powers” table created in part C. To make the implementation easier, we have provided a boilerplate for this exercise in the following file: **TablePartD.py**.

Read the values for the following row ids in order and the given attributes:

Id: "row1", Values for (hero, power, name, xp, color)

Id: "row19", Values for (hero, color)

Id: "row1", Values for (hero, name, color)

We will run **python3 TablePartD.py**. Following is an example of output lines, please use spaces between the words. Your output should contain three lines for this exercise.

```
hero: b'no', color: b'green'
```

Exercise E: Scan Data

In this exercise, you will scan the data in the populated “powers” table created in part C. To make the implementation easier, we have provided a boilerplate for this exercise in the following file: **TablePartE.py**.

We will run **python3 TablePartE.py**. The output should look similar to below:

```
root@24c5e229888:/MP3_py# python3 TablePartE.py
Found: b'row1', (b'professional:name': (b'batman', 1528643194884), b'custom:color': (b'black', 1528643194884), b'personal:hero': (b'yes', 1528643194884), b'professional:xp': (b'188', 1528643194884), b'personal:power': (b'fly', 1528643194884))
Found: b'row18', (b'professional:name': (b'gandalf', 1528643194151), b'custom:color': (b'grey', 1528643194151), b'personal:hero': (b'no', 1528643194151), b'professional:xp': (b'86', 1528643194151), b'personal:power': (b'think', 1528643194151))
Found: b'row11', (b'professional:name': (b'catman', 1528643194157), b'custom:color': (b'brown', 1528643194157), b'personal:hero': (b'no', 1528643194157), b'professional:xp': (b'84', 1528643194157), b'personal:power': (b'run', 1528643194157))
```

✓ Complete

